

Assignment 3

December 17, 2019

You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

1 Assignment 3 - Evaluation

In this assignment you will train several models and evaluate how effectively they predict instances of fraud using data based on [this dataset from Kaggle](#). Each row in `fraud_data.csv` corresponds to a credit card transaction. Features include confidential variables `V1` through `V28` as well as `Amount` which is the amount of the transaction. The target is stored in the `class` column, where a value of 1 corresponds to an instance of fraud and 0 corresponds to an instance of not fraud.

```
In [18]: import numpy as np
import pandas as pd
```

1.0.1 Question 1

Import the data from `fraud_data.csv`. What percentage of the observations in the dataset are instances of fraud?

This function should return a float between 0 and 1.

```
In [3]: # df = pd.read_csv('fraud_data.csv')
# bincount = np.bincount(df.Class)
# print(len(df))
# bincount
```

```
21693
```

```
Out[3]: array([21337,    356])
```

```
In [19]: def answer_one():

    df = pd.read_csv('fraud_data.csv')
    bincount = np.bincount(df.Class)
    return bincount[1]/len(df)
```

```
In [20]: # Use X_train, X_test, y_train, y_test for all of the following questions
         from sklearn.model_selection import train_test_split

         df = pd.read_csv('fraud_data.csv')

         X = df.iloc[:, :-1]
         y = df.iloc[:, -1]

         X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=0)
```

1.0.2 Question 2

Using `X_train`, `X_test`, `y_train`, and `y_test` (as defined above), train a dummy classifier that classifies everything as the majority class of the training data. What is the accuracy of this classifier? What is the recall?

This function should return a tuple with two floats, i.e. (accuracy score, recall score).

```
In [23]: def answer_two():
         from sklearn.dummy import DummyClassifier
         from sklearn.metrics import accuracy_score, recall_score

         dummy_clf = DummyClassifier(strategy = 'most_frequent').fit(X_train, y_train)
         predicted = dummy_clf.predict(X_test)

         return accuracy_score(y_test, predicted), recall_score(y_test, predicted)

In [24]: answer_two()

Out[24]: (0.98525073746312686, 0.0)
```

1.0.3 Question 3

Using `X_train`, `X_test`, `y_train`, `y_test` (as defined above), train a SVC classifier using the default parameters. What is the accuracy, recall, and precision of this classifier?

This function should return a tuple with three floats, i.e. (accuracy score, recall score, precision score).

```
In [27]: def answer_three():
         from sklearn.metrics import accuracy_score, recall_score, precision_score
         from sklearn.svm import SVC

         svm = SVC().fit(X_train, y_train)
         predicted = svm.predict(X_test)

         return accuracy_score(y_test, predicted), recall_score(y_test, predicted), precision_score(y_test, predicted)

In [28]: answer_three()

Out[28]: (0.99078171091445433, 0.375, 1.0)
```

1.0.4 Question 4

Using the SVC classifier with parameters {'C': 1e9, 'gamma': 1e-07}, what is the confusion matrix when using a threshold of -220 on the decision function. Use X_test and y_test.

This function should return a confusion matrix, a 2x2 numpy array with 4 integers.

```
In [8]: def answer_four():
        from sklearn.metrics import confusion_matrix
        from sklearn.svm import SVC

        svm = SVC(gamma=1e-07, C=1e9).fit(X_train, y_train)
        prediction = svm.decision_function(X_test) > -220
        confusion = confusion_matrix(y_test, prediction)

        return confusion
```

1.0.5 Question 5

Train a logistic regression classifier with default parameters using X_train and y_train.

For the logistic regression classifier, create a precision recall curve and a roc curve using y_test and the probability estimates for X_test (probability it is fraud).

Looking at the precision recall curve, what is the recall when the precision is 0.75?

Looking at the roc curve, what is the true positive rate when the false positive rate is 0.16?

This function should return a tuple with two floats, i.e. (recall, true positive rate).

```
In [10]: from sklearn.linear_model import LogisticRegression
        from sklearn.metrics import precision_recall_curve, roc_curve, auc

        # %matplotlib notebook
        # import seaborn as sns
        # import matplotlib.pyplot as plt
        # lr = LogisticRegression().fit(X_train, y_train)
        # lr_predicted = lr.predict(X_test)
        # precision, recall, thresholds = precision_recall_curve(y_test, lr_predicted)
        # fpr_lr, tpr_lr, _ = roc_curve(y_test, lr_predicted)

        # closest_zero = np.argmin(np.abs(thresholds))
        # closest_zero_p = precision[closest_zero]
        # closest_zero_r = recall[closest_zero]
        # plt.figure()
        # plt.xlim([0.0, 1.01])
        # plt.ylim([0.0, 1.01])
        # plt.plot(precision, recall, label='Precision-Recall Curve')
        # plt.plot(closest_zero_p, closest_zero_r, 'o', markersize = 12, fillstyle='none')
        # plt.xlabel('Precision', fontsize=16)
        # plt.ylabel('Recall', fontsize=16)
        # plt.axes().set_aspect('equal')
        # plt.show()
```

```
# roc_auc_lr = auc(fpr_lr, tpr_lr)
# plt.figure()
# plt.xlim([-0.01, 1.00])
# plt.ylim([-0.01, 1.01])
# plt.plot(fpr_lr, tpr_lr, lw=3, label='LogRegr ROC curve (area = {:.2f})')
# plt.xlabel('False Positive Rate', fontsize=16)
# plt.ylabel('True Positive Rate', fontsize=16)
# plt.title('ROC curve (1-of-10 digits classifier)', fontsize=16)
# plt.legend(loc='lower right', fontsize=13)
# plt.plot([0, 1], [0, 1], color='navy', lw=3, linestyle='--')
# plt.axes().set_aspect('equal')
# plt.show()
```

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

```
In [12]: def answer_five():
        return 0.83, 0.94
```

1.0.6 Question 6

Perform a grid search over the parameters listed below for a Logistic Regression classifier, using recall for scoring and the default 3-fold cross validation.

```
'penalty': ['l1', 'l2']
'C': [0.01, 0.1, 1, 10, 100]
```

From `.cv_results_`, create an array of the mean test scores of each parameter combination. i.e.

	l1	l2
0.01	?	?
0.1	?	?
1	?	?
10	?	?
100	?	?

This function should return a 5 by 2 numpy array with 10 floats.

Note: do not return a DataFrame, just the values denoted by '?' above in a numpy array. You might need to reshape your raw result to meet the format we are looking for.

```

In [13]: def answer_six():
          from sklearn.model_selection import GridSearchCV
          from sklearn.linear_model import LogisticRegression

          lr = LogisticRegression().fit(X_train, y_train)
          grid_values = {'penalty': ['l1', 'l2'], 'C': [0.01, 0.1, 1, 10, 100]}
          grid_clf_rec = GridSearchCV(lr, param_grid = grid_values, scoring = 'r2')
          grid_clf_rec.fit(X_train, y_train)
          return np.array(grid_clf_rec.cv_results_['mean_test_score']).reshape(5,2)

In [14]: answer_six()

Out[14]: array([[ 0.66666667,  0.76086957],
                 [ 0.80072464,  0.80434783],
                 [ 0.8115942 ,  0.8115942 ],
                 [ 0.80797101,  0.8115942 ],
                 [ 0.80797101,  0.80797101]])

In [15]: # Use the following function to help visualize results from the grid search
def GridSearch_Heatmap(scores):
    %matplotlib notebook
    import seaborn as sns
    import matplotlib.pyplot as plt
    plt.figure()
    sns.heatmap(scores.reshape(5,2), xticklabels=['l1','l2'], yticklabels=
    plt.yticks(rotation=0);

    #GridSearch_Heatmap(answer_six())

<IPython.core.display.Javascript object>

<IPython.core.display.HTML object>

In [ ]:

```